

Hierarchical Width-Based Planning and Learning

Miquel Junyent, Vicenç Gómez, Anders Jonsson

Universitat Pompeu Fabra, Barcelona, Spain



ICAPS
August 2021

IW(w) (Lipovetzky and Geffner 2012)

- Original algorithm: **Breadth-first search** (BrFS) method
- Requires states to be factored into **features** $\phi(s)$
- Prunes states that are not **novel for width w**
- A state is **novel** if it has a **feature tuple of size w** that is **new in the search**
- **Complexity exponential in w** , but independent of $|S|$
- Most classical planning benchmarks present a **low width** when the **goal is a single atom**.

# Domains	# Inst.	Inst. IW(1)	Inst. IW(2)
37	37,921	37.0%	88.2%

Summary of Table 1 (Lipovetzky and Geffner 2012)

IW(w) (Lipovetzky and Geffner 2012)

- Original algorithm: **Breadth-first search** (BrFS) method
- Requires states to be factored into **features** $\phi(s)$
- Prunes states that are not **novel for width w**
- A state is **novel** if it has a **feature tuple of size w** that is **new in the search**
- **Complexity exponential in w** , but independent of $|S|$
- Most classical planning benchmarks present a **low width** when the **goal is a single atom**.
- In practice:
 - Problems have **higher width** (no single-atom goal tasks)
 - IW(w) is mostly used with **$w=1$** due to computational constraints

Complexity of IW(w)

- Let $N(n, d, w)$ denote the **maximum amount of novel nodes** that IW(w) generates in a problem with $n = |F|$ features of domain size $d = |D|$
- Our result is based on two basic premises:
 - A feature **has one value** at a time
 - A feature value **appears in several tuples** simultaneously
- Recursive formula

Complexity of IW(w)

- Let $N(n, d, w)$ denote the **maximum amount of novel nodes** that IW(w) generates in a problem with $n = |F|$ features of domain size $d = |D|$
- Our result is based on two basic premises:
 - A feature **has one value** at a time
 - A feature **appears in several tuples** simultaneously
- Recursive formula

$$N(n, d, 0) = 1, \longrightarrow \text{Only the initial state is novel}$$

$$w = 0$$

Complexity of IW(w)

- Let $N(n, d, w)$ denote the **maximum amount of novel nodes** that IW(w) generates in a problem with $n = |F|$ features of domain size $d = |D|$
- Our result is based on two basic premises:
 - A feature **has one value** at a time
 - A feature **appears in several tuples** simultaneously
- Recursive formula

$$N(n, d, 0) = 1,$$

$$N(n, d, n) = d^n, \longrightarrow \text{All states are novel}$$

$$n = w$$

Complexity of IW(w)

- Let $N(n, d, w)$ denote the **maximum amount of novel nodes** that IW(w) generates in a problem with $n = |F|$ features of domain size $d = |D|$
- Our result is based on two basic premises:
 - A feature **has one value** at a time
 - A feature **appears in several tuples** simultaneously
- Recursive formula

$$N(n, d, 0) = 1,$$

$$N(n, d, n) = d^n,$$

$$N(n, d, w) = \underbrace{(d - 1)N(n - 1, d, w - 1)}_{\text{States novel due to one feature } f} + \underbrace{N(n - 1, d, w)}_{\text{States novel due to other features different than } f}.$$

States novel due
to **one feature** f

States novel due
to **other features**
different than f

Complexity of IW(w)

- Let $N(n, d, w)$ denote the **maximum amount of novel nodes** that IW(w) generates in a problem with $n = |F|$ features of domain size $d = |D|$
- Our result is based on two basic premises:
 - A feature **has one value** at a time
 - A feature **appears in several tuples** simultaneously
- General formula, for $0 \leq w < n$:

$$N(n, d, w) = \sum_{k=0}^w \left[\binom{n-1-k}{w-k} d^k (d-1)^{w-k} \right].$$

A hierarchical approach to blind search

- Blind search methods require two components:
 - **Successor function:** given a state and an action, returns a successor state (e.g., simulator)
 - **Stopping condition:** tells us when to stop the search (e.g., goal is met)

Example BrFS:

```
Q = Queue(root)
While Q not empty:
    s = PopFirst(Q)
    For each action a:
        x = GenerateSuccessor(s, a)
        Append(Q, x)
        If ShouldStop(x):
            return
```

A hierarchical approach to blind search

- Our **hierarchical approach**:
 - Considers two sets of features F_h and F_ℓ

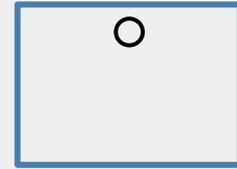
A state is represented by its feature vector $\phi_\ell(s)$

○

A hierarchical approach to blind search

- Our **hierarchical approach**:
 - Considers two sets of features F_h and F_ℓ

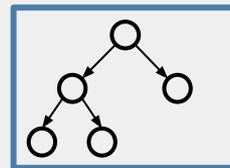
It also belongs to a **high-level state**, represented by the feature mapping $\phi_h(s)$



A hierarchical approach to blind search

- Our **hierarchical approach**:
 - Considers two sets of features F_h and F_ℓ
 - Modifies:
 - **High-level successor function**: each call to this function triggers a low-level search

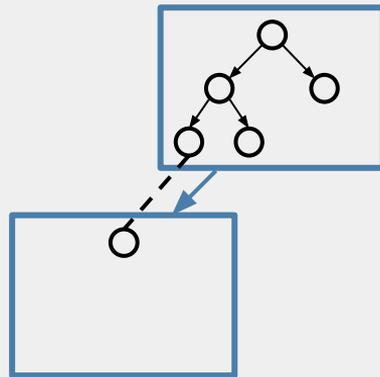
Each **high-level state** contains a **low-level tree**, where states share the same $\phi_h(s)$



A hierarchical approach to blind search

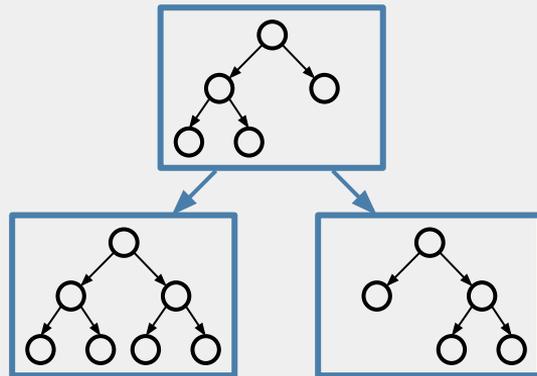
- Our **hierarchical approach**:
 - Considers two sets of features F_h and F_ℓ
 - Modifies:
 - **High-level successor function**: each call to this function triggers a low-level search
 - **Low-level stopping condition**: stops the low-level search when a state s that maps to a different $\phi_h(s)$ is found

When a state maps to a **different** $\phi_h(s)$, the low-level search is stopped and a **new high-level node** is created



A hierarchical approach to blind search

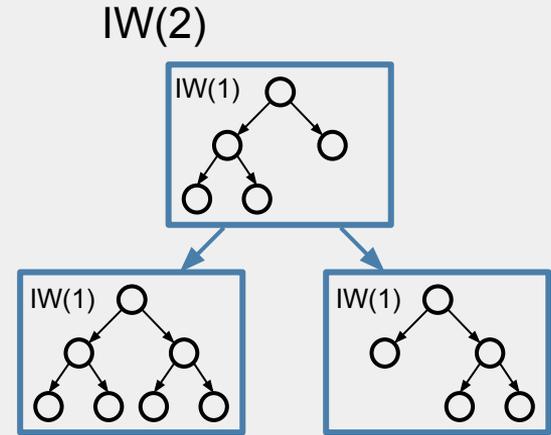
- Our **hierarchical approach**:
 - Considers two sets of features F_h and F_ℓ
 - Modifies:
 - **High-level successor function**: each call to this function triggers a low-level search
 - **Low-level stopping condition**: stops the low-level search when a state s that maps to a different $\phi_h(s)$ is found
 - We can pause and **resume low-level searches**
 - Allows for **many levels of abstraction**
 - Accepts **different planners** at each level



Hierarchical IW

- We can use **IW(w)** at the different levels
- For instance:
 - IW(2) at high-level
 - IW(1) at low-level

} HIW(2, 1)



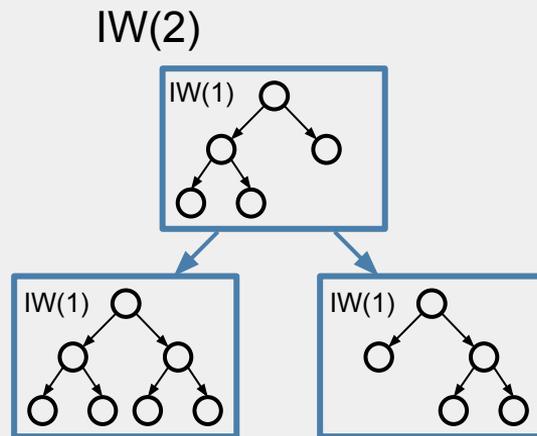
Hierarchical IW

- We can use **IW(w)** at the different levels
- For instance:
 - IW(2) at high-level
 - IW(1) at low-level
- In general: $HIW(w_h, w_\ell)$
- HIW can solve problems of width $w_h + w_\ell$



Features:

- 1-D position
- Having the key



Hierarchical IW

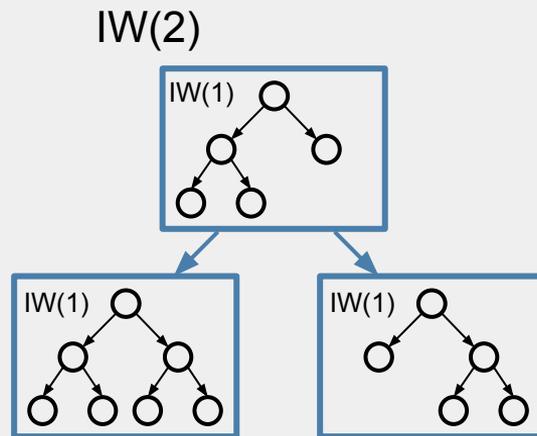
- We can use **IW(w)** at the different levels
 - For instance:
 - IW(2) at high-level
 - IW(1) at low-level
- } HIW(2, 1)
- In general: $\text{HIW}(w_h, w_\ell)$
 - HIW can solve problems of width $w_h + w_\ell$



Features:

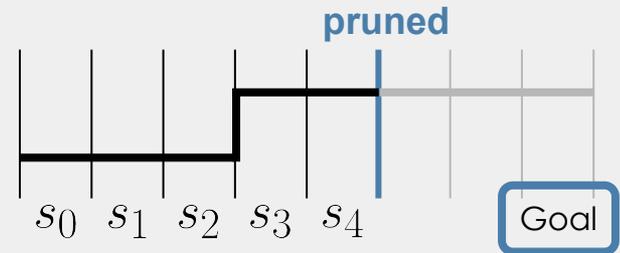
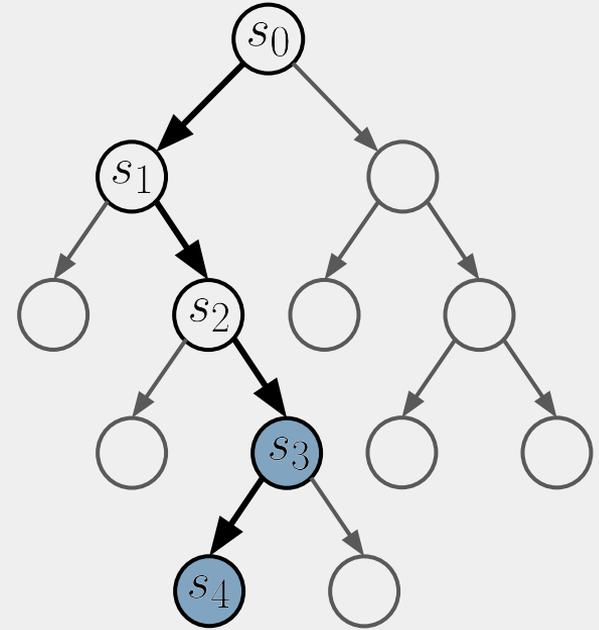
- 1-D position (**low level**)
- Having the key (**high level**)

This problem has **width 2**,
but it can be solved by
HIW(1,1)



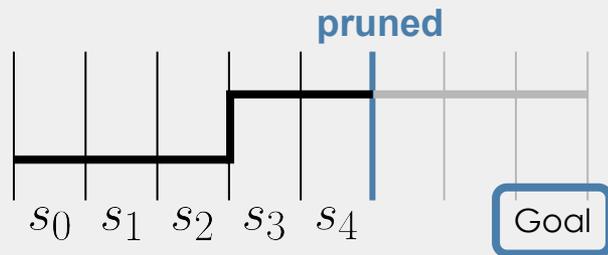
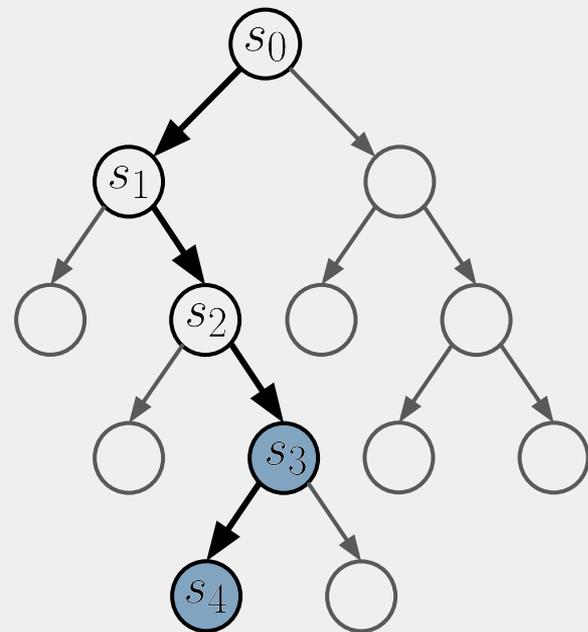
Incremental HIW

- **Hypothesis:** features that **only change once** in a branch before being pruned are good candidates



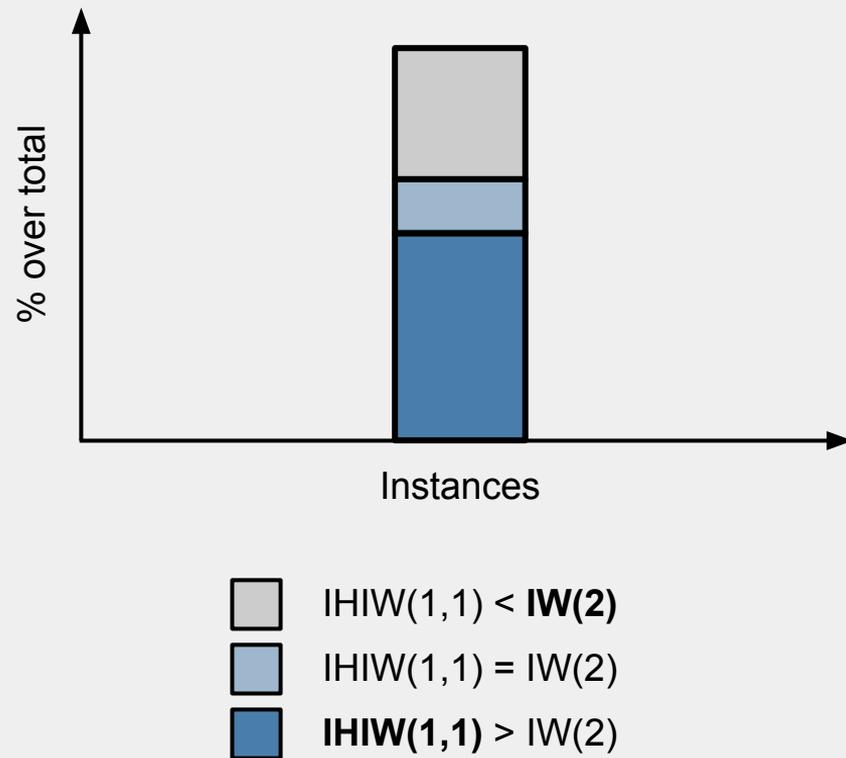
Incremental HIW

- **Hypothesis:** features that **only change once** in a branch before being pruned are good candidates
- **Incremental HIW(1,1):**
 - Iteratively **run HIW(1,1)**
 - **Add one feature** to F_h at each iteration
 - **Discover new features** when necessary
 - **Reuse** the search **tree** among iterations



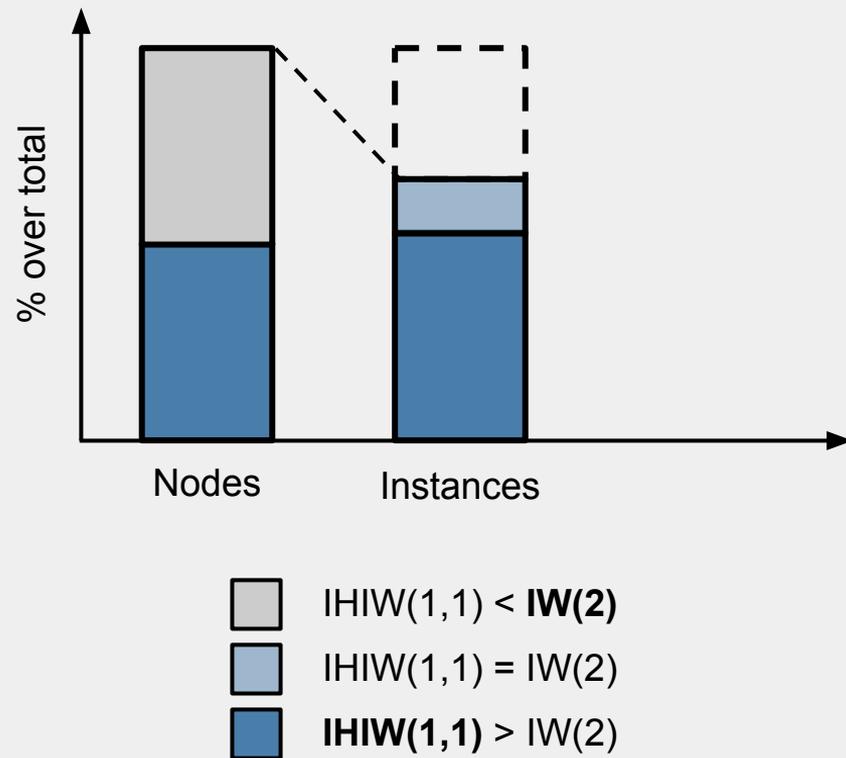
Results in classical planning

- Single goal instances
- Budget of 10K nodes
- We report:
 - Solved instances (%)
 - Avg. nodes (solved)
 - Avg. time (solved)
- **IHIW > IW(1) in 31/36 domains**
- **IHIW \geq IW(2) in 24/36 domains**



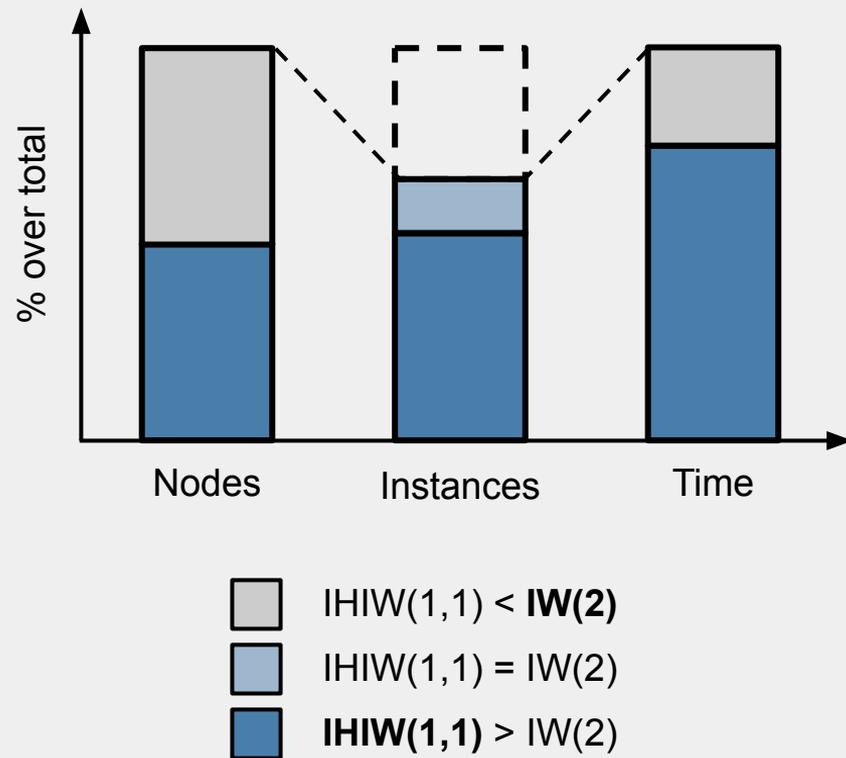
Results in classical planning

- Single goal instances
- Budget of 10K nodes
- We report:
 - Solved instances (%)
 - Avg. nodes (solved)
 - Avg. time (solved)
- **IHIW > IW(1) in all but 5 domains**
- **IHIW ≥ IW(2) in 24/36 domains**
 - Uses **less nodes in 12/24**



Results in classical planning

- Single goal instances
- Budget of 10K nodes
- We report:
 - Solved instances (%)
 - Avg. nodes (solved)
 - Avg. time (solved)
- **IHIW > IW(1) in all but 5 domains**
- **IHIW ≥ IW(2) in 24/36 domains**
 - Uses **less nodes** in 12/24
 - Solves it **faster** in 18/24



π -HIW

- Integrating HIW with a **policy learning** scheme



π -HIW

- Integrating HIW with a **policy learning** scheme
 - **High-level planner: Count-Based Rollout IW**
 - Selects high-level nodes according to $p \propto \exp(1/\tau(c+1))$
 - Prunes nodes using a mapping from novel tuples to unpruned nodes



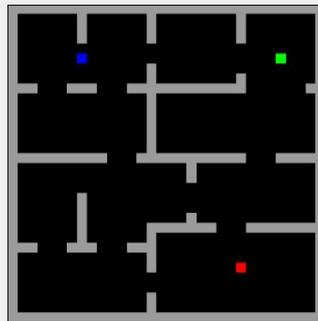
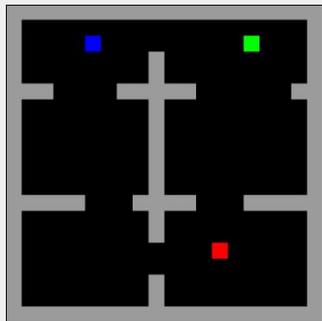
π -HIW

- Integrating HIW with a **policy learning** scheme
 - **High-level planner: Count-Based Rollout IW**
 - Selects high-level nodes according to $p \propto \exp(1/\tau(c+1))$
 - Prunes nodes using a mapping from novel tuples to unpruned nodes
 - **Low-level planner: π -IW modified**
 - Tree counts for tie-breaking (when the reward is sparse)
 - Value function



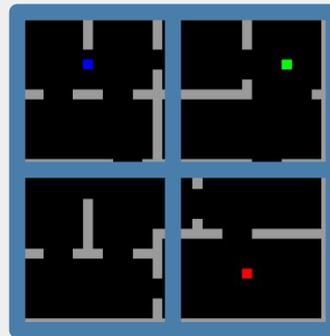
Results in gridworld

- Sparse reward tasks. The episode terminates:
 - when the agent ■ picks the key ■ and reaches the door ■ ($r = +1$)
 - when hitting a wall ($r = -1$)
 - after 200 / 500 steps ($r = 0$)



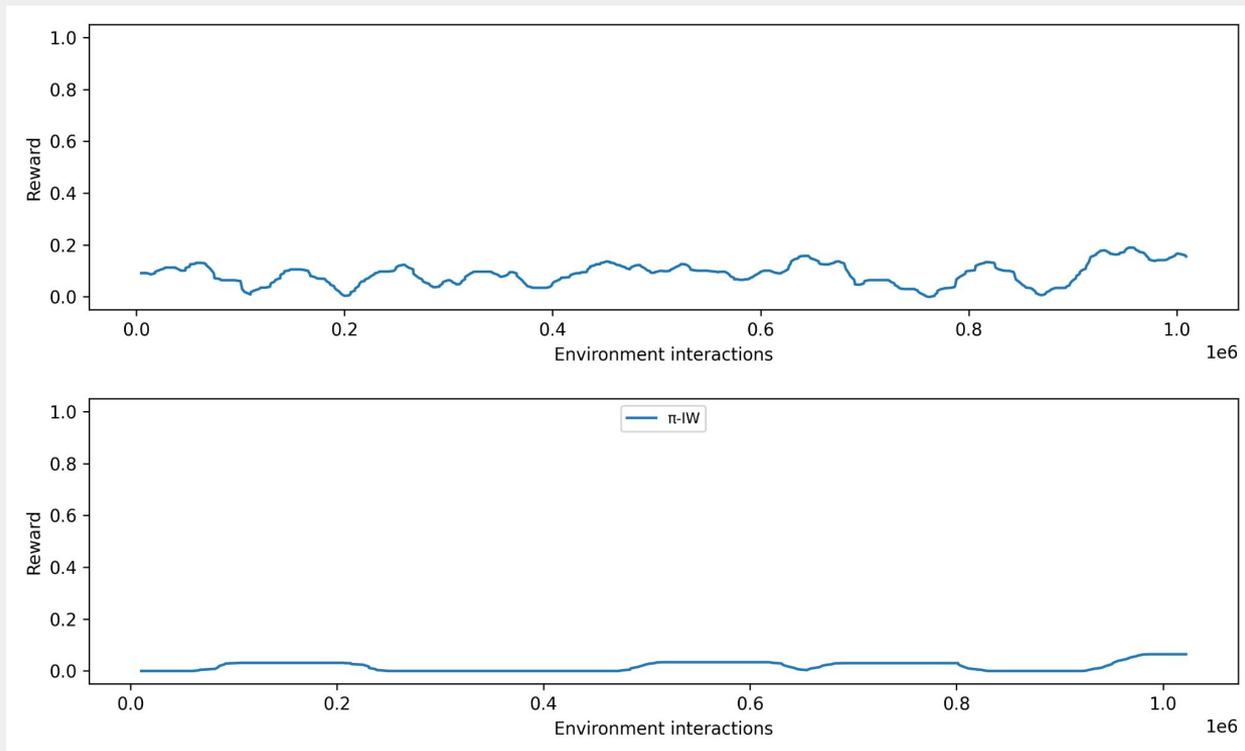
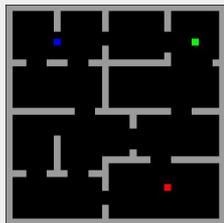
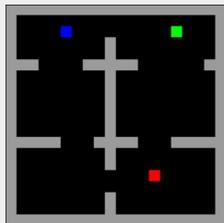
Results in gridworld

- Sparse reward tasks. The episode terminates:
 - when the agent ■ picks the key ■ and reaches the door ■ ($r = +1$)
 - when hitting a wall ($r = -1$)
 - after 200 / 500 steps ($r = 0$)
- Features:
 - Neural network activations (low level)
 - Downsampling (high level)



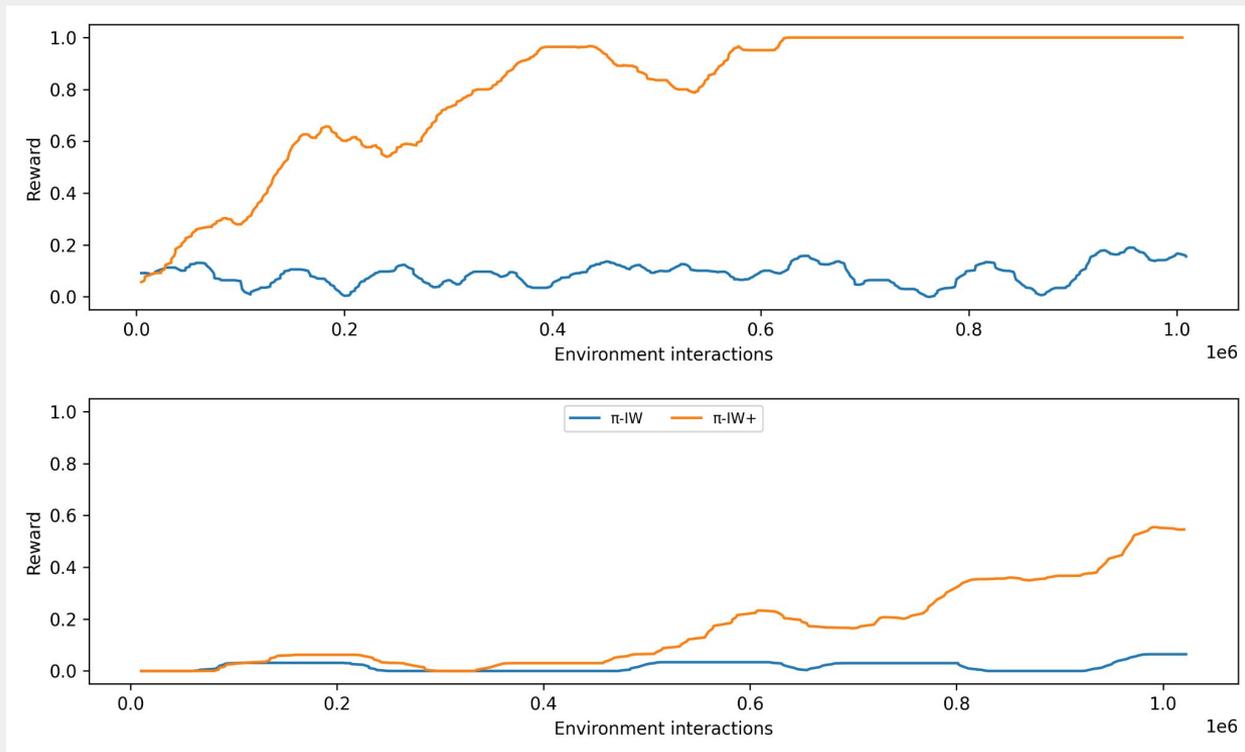
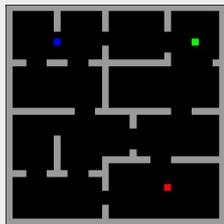
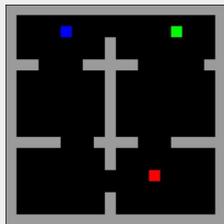
Example: 2x2 tiles

Results in gridworld



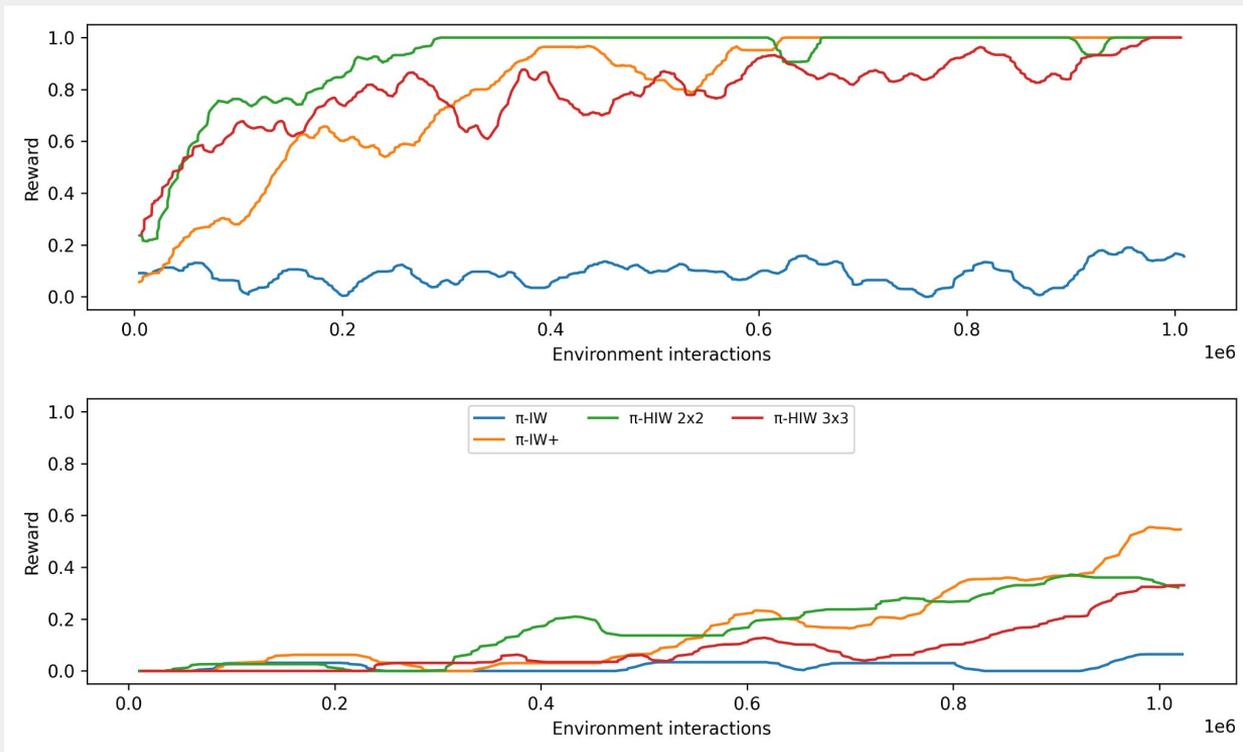
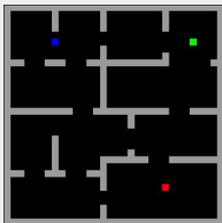
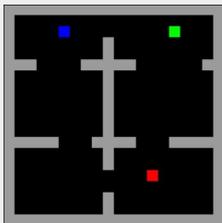
Features: neural network activations

Results in gridworld



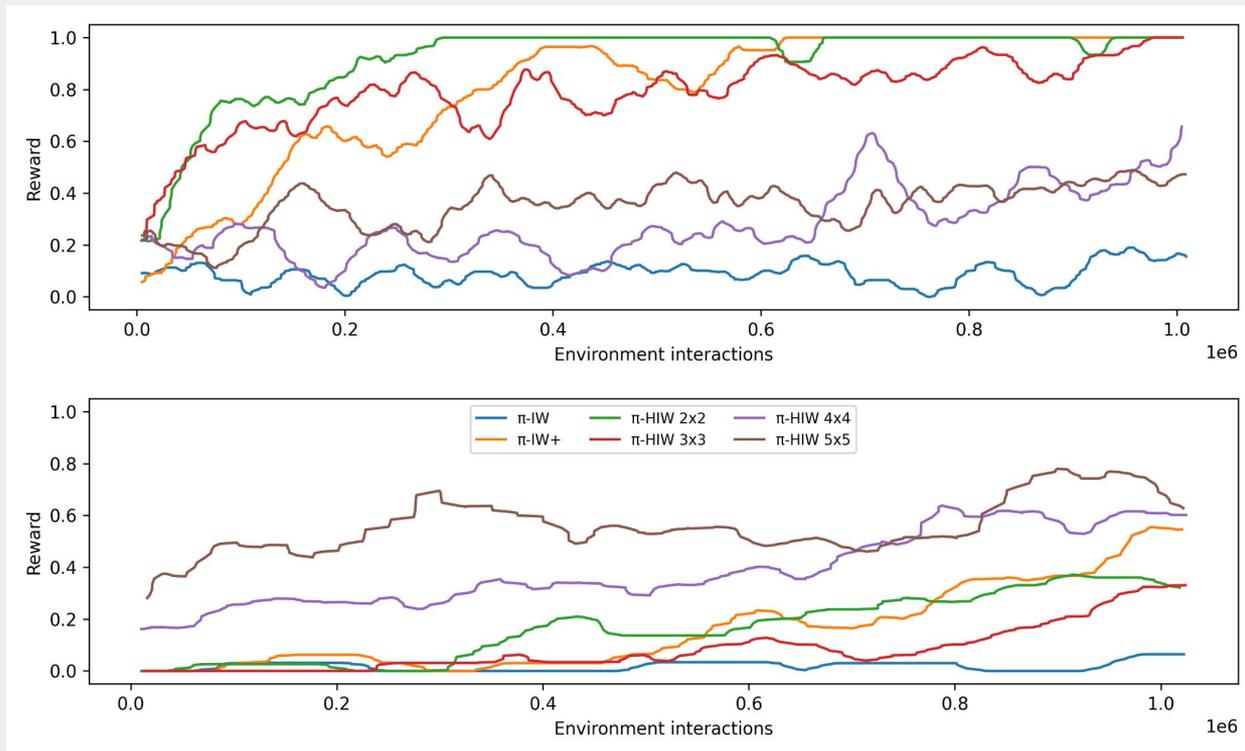
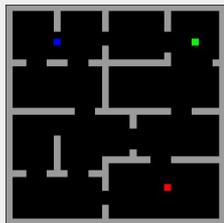
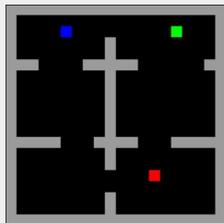
Features: neural network activations

Results in gridworld



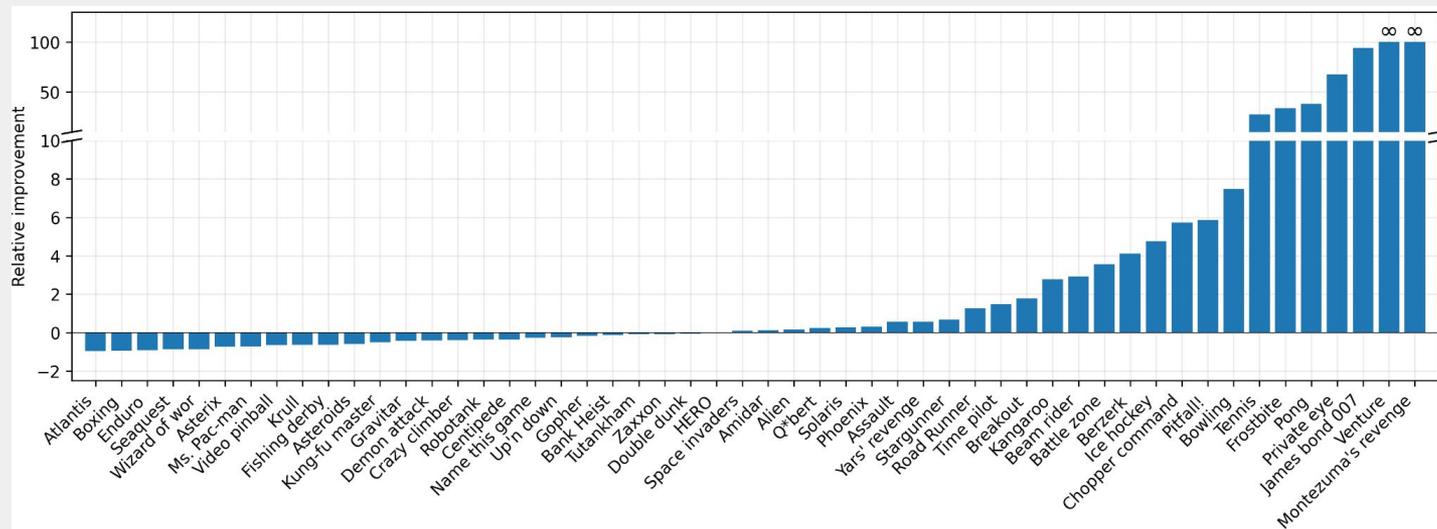
Features: neural network activations (**low level**)
downsampling (**high level**)

Results in gridworld

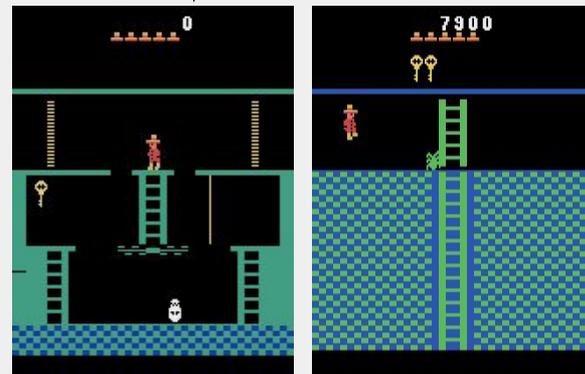
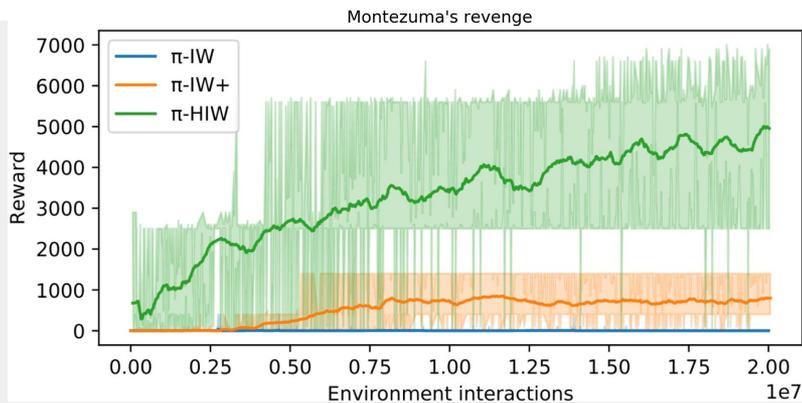
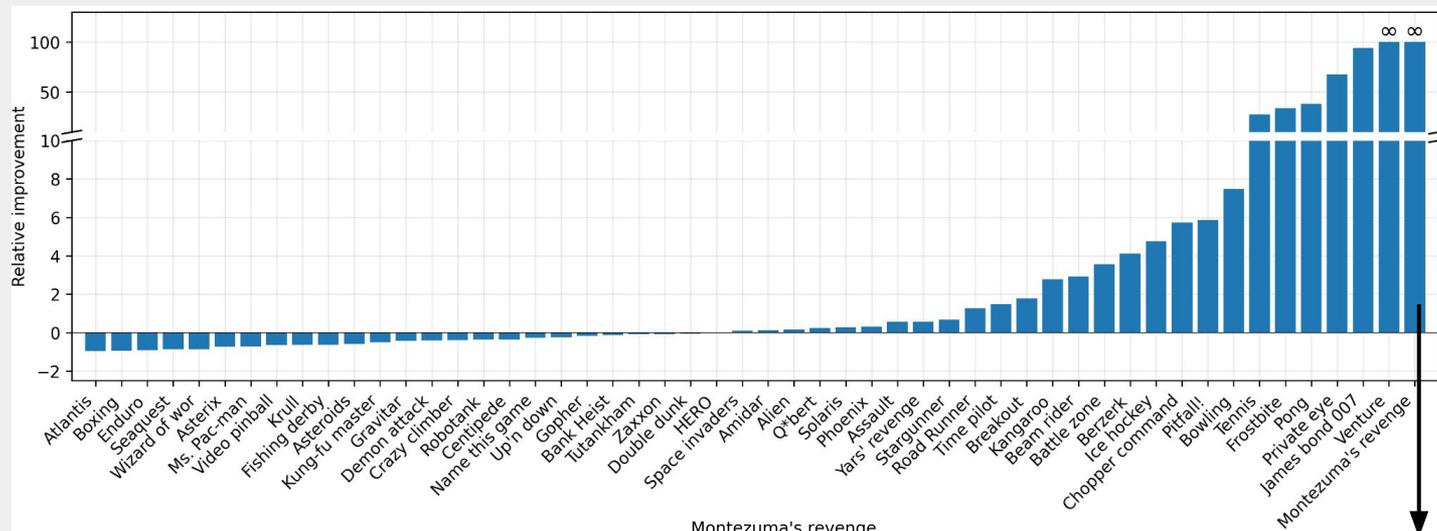


Features: neural network activations (**low level**)
downsampling (**high level**)

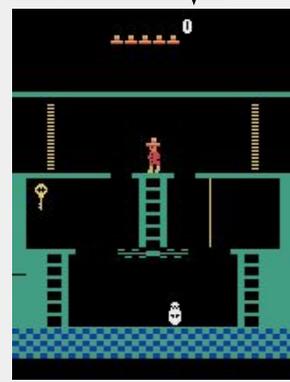
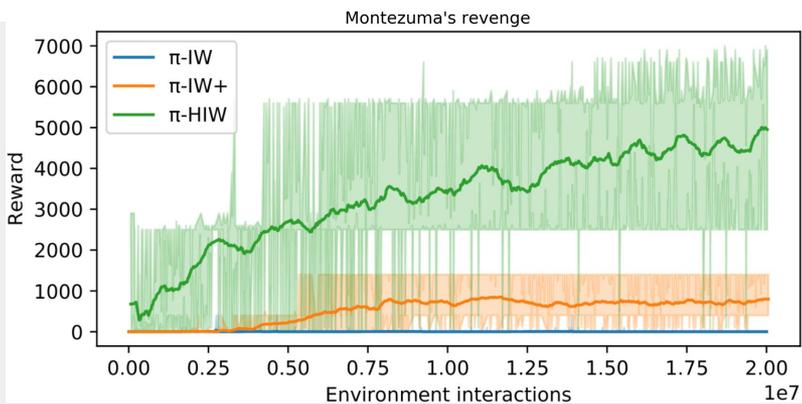
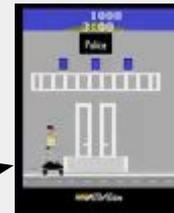
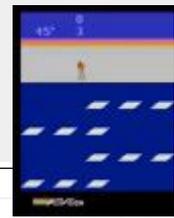
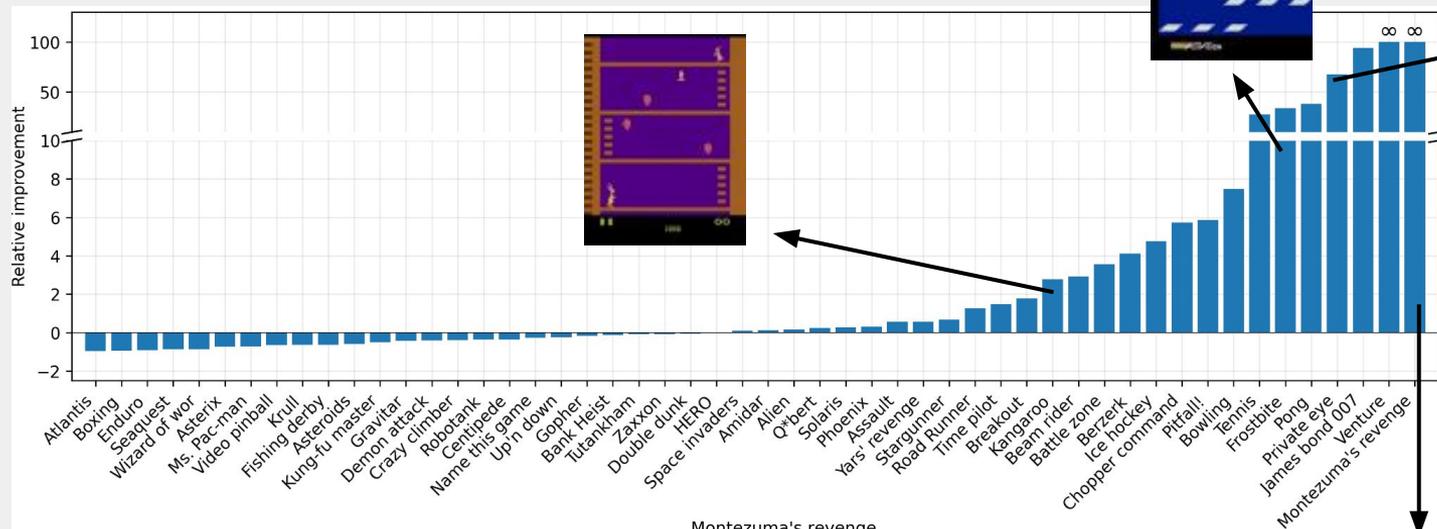
Results in Atari games



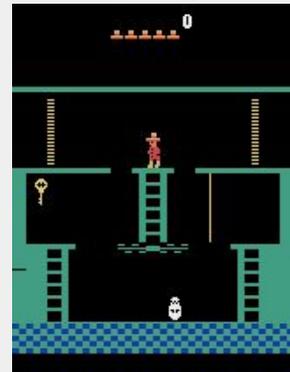
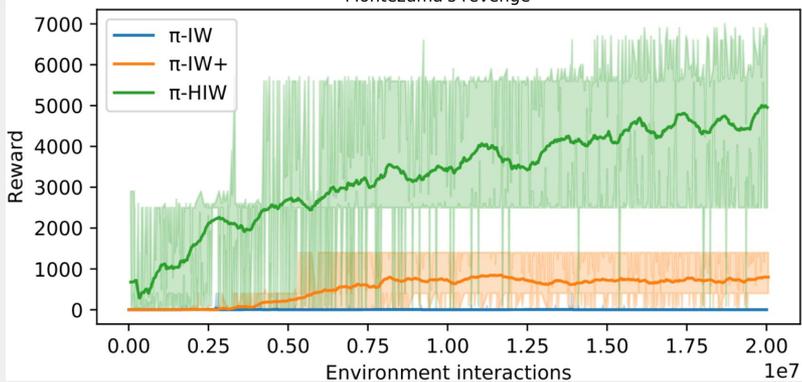
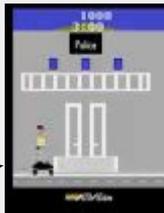
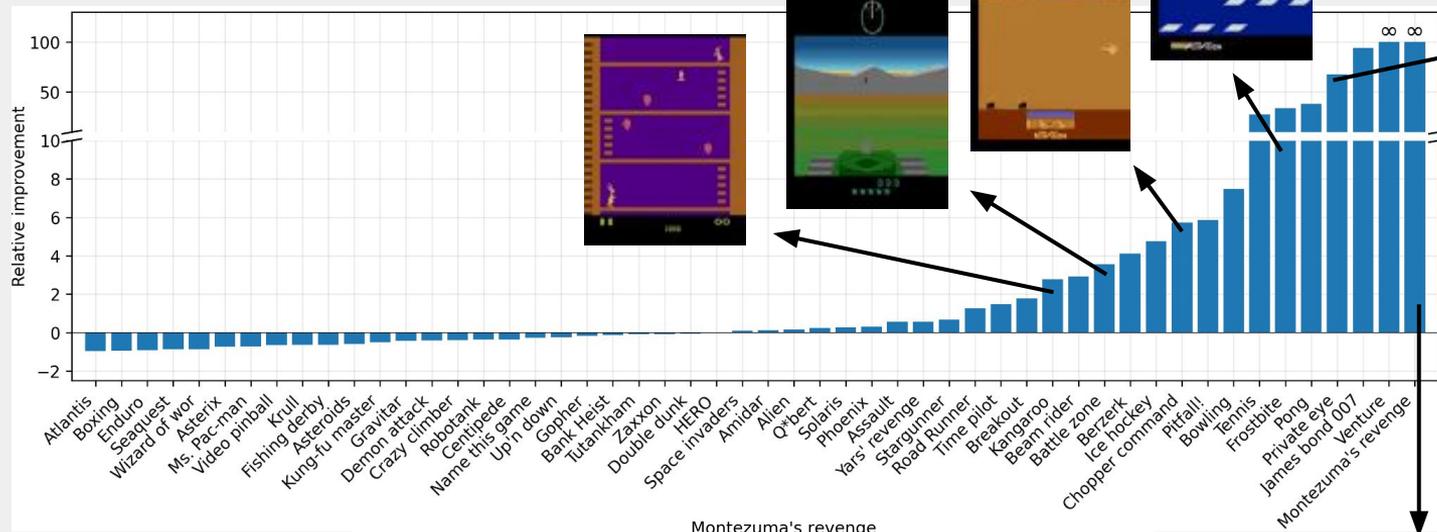
Results in Atari games



Results in Atari games



Results in Atari games



Conclusions

- **Tighter bound** for $IW(w)$
- Simple method for **hierarchical blind search**
- **Hierarchical IW:**
 - Can solve problems of width $w_h + w_\ell$
 - Incremental HIW: **discovering high-level features**
 - Experiments: **IHIW > IW(2)** in 24/36 domains (> instances, < nodes, < time)
- **π -HIW:**
 - **Count-based Rollout IW** (high-level planner)
 - **Improvements to π -IW** (low-level planner):
 - Value
 - Counts
 - Experiments: gridworld and Atari games

Thanks!

<https://github.com/aig-upf/hierarchical-IW>